

Homework -2

MULTIPLE REGRESSION

```
In [99]: import numpy as np
import pandas as pd
from scipy import stats

import matplotlib as ml
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

```
In [84]: data_df = pd.read_csv('data.csv')
```

```
In [85]: data_df.head()
```

Out[85]:

| | Hour | CO | Traffic | Wind |
|---|------|-----|---------|------|
| 0 | 1 | 2.4 | 50 | -0.2 |
| 1 | 2 | 1.7 | 26 | 0.0 |
| 2 | 3 | 1.4 | 16 | 0.0 |
| 3 | 4 | 1.2 | 10 | 0.0 |
| 4 | 5 | 1.2 | 12 | 0.1 |

1. Find the correlation and plot the scatter plot of (Y and H), (Y and T), (Y and W)

```
In [86]: Y = data_df.CO  
H = data_df.Hour  
T = data_df.Traffic  
W = data_df.Wind
```

```
In [129]: print np.corrcoef(H,Y)  
print "\nHere Correlation between H and Y is 43.83%"  
  
print np.corrcoef(T,Y)  
print "\nHere Correlation between T and Y is 96.26%"  
  
print np.corrcoef(W,Y)  
print "\nHere Correlation between W and Y is 70.97%"
```

```
[[ 1.          0.42835141]
 [ 0.42835141  1.          ]]
```

Here Correlation between H and Y is 43.83%

```
[[ 1.          0.96266546]
 [ 0.96266546  1.          ]]
```


Here Correlation between T and Y is 96.26%

```
[[ 1.          0.70974865]
 [ 0.70974865  1.          ]]
```

Here Correlation between W and Y is 70.97%

```
In [145]: fig, axs = plt.subplots(1,3)
data_df.plot(kind = "scatter", x = 'Hour', y = 'CO', ax = axs[0], figsize=(18,6))
data_df.plot(kind = 'scatter', x = 'Traffic', y = 'CO', ax = axs[1])
data_df.plot(kind = 'scatter', x = 'Wind', y = 'CO', ax = axs[2])
```

```
Out[145]: <matplotlib.axes._subplots.AxesSubplot at 0x124081110>
```



2.

In [168]:


```
slope1, intercept1, r_value1, p_value1, std_err1 = stats.linregress(H,Y)
slope2, intercept2, r_value2, p_value2, std_err2 = stats.linregress(T,Y)
slope3, intercept3, r_value3, p_value3, std_err3 = stats.linregress(W,Y)
```

```
print "R^2 for H and Y is: ", r_value1**2
print "R^2 for T and Y is: ", r_value2**2
print "R^2 for W and Y is: ", r_value3**2
```

```
R^2 for H and Y is: 0.183484932457
R^2 for T and Y is: 0.92672478996
R^2 for W and Y is: 0.503743142976
```

```
In [144]: g = sns.PairGrid(data_df, y_vars=["CO"], x_vars=["Hour", "Traffic", "Wind"],
size=6)
g.map(sns.regplot)
```

Out[144]: <seaborn.axisgrid.PairGrid at 0x123553390>



3.

```
In [125]: print np.corrcoef(T,H)
print 'Correlation Between T and H is 42.85%'
print np.corrcoef(W,T)
print 'Correlation Between W and T is 61.34%'
print np.corrcoef(H,W)
print 'Correlation Between H and W is 42.26%'
```

```
[[ 1.          0.42852285]
 [ 0.42852285  1.          ]]
Correlation Between T and H is 42.85%
[[ 1.          0.61344359]
 [ 0.61344359  1.          ]]
Correlation Between W and T is 61.34%
[[ 1.          0.42263809]
 [ 0.42263809  1.          ]]
Correlation Between H and W is 42.26%
```

4.

```
In [201]: # Y = Y , X = T,H
Y = data_df['CO']
X1 = data_df[['Traffic', 'Hour']]
X2 = data_df[['Wind', 'Traffic']]
X3 = data_df[['Hour', 'Wind']]
X = data_df[['Hour', 'Traffic', 'Wind']]
```

Traffic and Hour

```
In [178]: import statsmodels.formula.api as smf
lm1 = smf.ols(formula='CO ~ Hour+Traffic', data=data_df).fit()
```

```
In [185]: print "Coefficients of the model: "
print lm1.params
```

Coefficients of the model:

```
Intercept    1.175000
Hour         0.005815
Traffic      0.020647
dtype: float64
```

```
In [186]: print "Confidence Interval: "
          print lm1.conf_int()
```

Confidence Interval:

| | 0 | 1 |
|-----------|-----------|----------|
| Intercept | 0.593308 | 1.756692 |
| Hour | -0.034879 | 0.046509 |
| Traffic | 0.017712 | 0.023583 |

```
In [187]: print "R-squared: "
          print lm1.rsquared
```

R-squared:

0.927031639809

```
In [188]: print "Predictions of CO using model i.e due to regression: "
          print lm1.predict(X1)
```

Predictions of CO using model i.e due to regression:

| | | | | | |
|--------------|------------|------------|------------|------------|-------------|
| [2.21318299 | 1.72346137 | 1.52280278 | 1.40473362 | 1.45184334 | 2.05643174 |
| 4.45734036 | 6.92019105 | 7.0498902 | 6.22981084 | 5.36843677 | 5.08518874 |
| 4.94647222 | 4.93163986 | 5.45363882 | 6.71894271 | 7.24094167 | 7.63905648 |
| 6.79832977 | 5.21429814 | 3.87803481 | 3.78061301 | 3.70383857 | 3.11088016] |

wind and aTraffic

```
In [189]: lm2 = smf.ols(formula='CO ~ Wind+Traffic', data=data_df).fit()
```

```
In [190]: print "Coefficients of the model: "  
print lm2.params
```

```
Coefficients of the model:  
Intercept    1.274461  
Wind         0.174747  
Traffic      0.018290  
dtype: float64
```

```
In [191]: print "Confidence Interval: "  
print lm2.conf_int()
```

```
Confidence Interval:  
                0          1  
Intercept  0.862413  1.686510  
Wind       0.056697  0.292797  
Traffic    0.015497  0.021084
```

```
In [192]: print "R-squared: "  
print lm2.rsquared
```

```
R-squared:  
0.949509414529
```

```
In [194]: print "Predictions of CO using model i.e due to regression: "  
print lm2.predict(X2)
```


Predictions of CO using model i.e due to regression:

```
[ 2.15403275  1.75001201  1.56710781  1.45736529  1.5114208  2.00689364
 4.12858237  6.28766769  6.4673089  5.87548946  5.33446254  5.34051672
 5.35228115  5.4737881  6.01842196  7.13413758  7.53897407  7.76416936
 6.82204075  5.18650768  3.805409  3.57415953  3.46604851  2.88320231]
```

Hour and Wind

```
In [195]: lm3 = smf.ols(formula='CO ~ Hour+Wind', data=data_df).fit()
```

```
In [196]: print "Coefficients of the model: "
          print lm3.params
```

Coefficients of the model:

Intercept 2.776934

Hour 0.046881

Wind 0.588501

dtype: float64

```
In [198]: print "Confidence Interval: "
          print lm3.conf_int()
```

Confidence Interval:

| | 0 | 1 |
|-----------|-----------|----------|
| Intercept | 1.432270 | 4.121597 |
| Hour | -0.056757 | 0.150520 |
| Wind | 0.272591 | 0.904411 |

```
In [199]: print "R-squared: "
          print lm3.rsquared
```

R-squared:
0.523810181436

```
In [200]: print "Predictions of CO using model i.e due to regression: "  
print lm3.predict(X3)
```

Predictions of CO using model i.e due to regression:

```
[ 2.70611473  2.87069625  2.91757757  2.96445888  3.0701903   2.99937142  
 3.04625274  3.03428395  3.31656567  3.83424779  4.64618041  5.57581323  
 6.09349535  6.61117747  6.95230929  6.99919061  6.86952163  6.50445224  
 5.90398246  5.18581247  4.58534268  4.1614232   4.09060432  3.96093533]
```

5

```
In [202]: lm = smf.ols(formula='CO ~ Hour+Traffic+Wind', data=data_df).fit()
```

```
In [203]: print "Coefficients of the model: "  
print lm.params
```

Coefficients of the model:

```
Intercept    1.318967  
Hour         -0.005689  
Traffic       0.018402  
Wind         0.179189  
dtype: float64
```

```
In [204]: print "Confidence Interval: "  
print lm.conf_int()
```

Confidence Interval:

| | 0 | 1 |
|-----------|-----------|----------|
| Intercept | 0.813076 | 1.824858 |
| Hour | -0.041289 | 0.029911 |
| Traffic | 0.015455 | 0.021349 |
| Wind | 0.055038 | 0.303340 |

```
In [205]: print "R-squared: "  
print lm.rsquared
```

R-squared:
0.949788420165

```
In [206]: print "Predictions of CO using model i.e due to regression: "  
print lm.predict(X)
```

Predictions of CO using model i.e due to regression:

| | | | | | |
|--------------|------------|------------|------------|------------|-------------|
| [2.19752889 | 1.78603442 | 1.59632719 | 1.48022712 | 1.52926028 | 2.02138505 |
| 4.15030336 | 6.31650814 | 6.49290518 | 5.89449551 | 5.34887679 | 5.35434601 |
| 5.36319541 | 5.48245555 | 6.02640551 | 7.14322537 | 7.54382407 | 7.76233641 |
| 6.80506571 | 5.14949276 | 3.7505791 | 3.50952957 | 3.39439527 | 2.80129731] |

```
In [208]: import os  
os.environ
```

Out[208]: {'TERM_PROGRAM_VERSION': '361.1', 'LOGNAME': 'dhrumindesai', 'USER': 'dhrumindesai', 'PATH': '/Users/dhrumindesai/anaconda/bin:/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin', 'HOME': '/Users/dhrumindesai', 'TERM_PROGRAM': 'Apple_Terminal', 'LANG': 'en_US.UTF-8', 'TERM': 'xterm-color', 'Apple_PubSub_Socket_Render': '/private/tmp/com.apple.launchd.h5sGFTrCQH/Render', 'SHLVL': '1', 'XPC_FLAGS': '0x0', 'TMPDIR': '/var/folders/6j/hxy_rx8d2713_pfzyb7l7mp00000gn/T/', 'TERM_SESSION_ID': '369B75E4-025E-435F-A3FD-78111711F9C9', 'XPC_SERVICE_NAME': '0', 'JPY_PARENT_PID': '568', 'SSH_AUTH_SOCK': '/private/tmp/com.apple.launchd.lJ3gWUSu6u/Listeners', 'SHELL': '/bin/bash', 'GIT_PAGER': 'cat', '_': '/Users/dhrumindesai/anaconda/bin/jupyter', 'OLDPWD': '/Users/dhrumindesai/desktop/Stevens/BIA-652', 'CLICOLOR': '1', '__CF_USER_TEXT_ENCODING': '0x1F5:0x0:0x0', 'PWD': '/Users/dhrumindesai/desktop/Stevens/BIA-652/Homework-2', 'PAGER': 'cat'}

In []: